

Case study on Numpy and Pandas

November 22, 2024

1 NumPy Problems

```
[114]: import numpy as np
        from numpy.linalg import norm
        import csv
```

```
[115]: with open('auto-mpg.csv', 'r') as file:
        reader = csv.reader(file)
        data = list(reader)
```

```
[116]: # Headers
        data[0]
```

```
[116]: ['mpg',
        'cylinders',
        'displacement',
        'horsepower',
        'weight',
        'acceleration',
        'model year',
        'origin',
        'car name']
```

1.1 1. Basic Array Operations

- Convert the mpg column into a NumPy array and calculate: The mean, median, and standard deviation of mpg.
- The number of cars with mpg greater than 25.

```
[117]: mpg_data = np.array([data[i][0] for i in range(1, len(data))]).astype('float64')
        mpg_data
```

```
[117]: array([18. , 15. , 18. , 16. , 17. , 15. , 14. , 14. , 14. , 15. , 15. ,
        14. , 15. , 14. , 24. , 22. , 18. , 21. , 27. , 26. , 25. , 24. ,
        25. , 26. , 21. , 10. , 10. , 11. , 9. , 27. , 28. , 25. , 25. ,
        19. , 16. , 17. , 19. , 18. , 14. , 14. , 14. , 14. , 12. , 13. ,
        13. , 18. , 22. , 19. , 18. , 23. , 28. , 30. , 30. , 31. , 35. ,
        27. , 26. , 24. , 25. , 23. , 20. , 21. , 13. , 14. , 15. , 14. ,
```

```

17. , 11. , 13. , 12. , 13. , 19. , 15. , 13. , 13. , 14. , 18. ,
22. , 21. , 26. , 22. , 28. , 23. , 28. , 27. , 13. , 14. , 13. ,
14. , 15. , 12. , 13. , 13. , 14. , 13. , 12. , 13. , 18. , 16. ,
18. , 18. , 23. , 26. , 11. , 12. , 13. , 12. , 18. , 20. , 21. ,
22. , 18. , 19. , 21. , 26. , 15. , 16. , 29. , 24. , 20. , 19. ,
15. , 24. , 20. , 11. , 20. , 21. , 19. , 15. , 31. , 26. , 32. ,
25. , 16. , 16. , 18. , 16. , 13. , 14. , 14. , 14. , 29. , 26. ,
26. , 31. , 32. , 28. , 24. , 26. , 24. , 26. , 31. , 19. , 18. ,
15. , 15. , 16. , 15. , 16. , 14. , 17. , 16. , 15. , 18. , 21. ,
20. , 13. , 29. , 23. , 20. , 23. , 24. , 25. , 24. , 18. , 29. ,
19. , 23. , 23. , 22. , 25. , 33. , 28. , 25. , 25. , 26. , 27. ,
17.5, 16. , 15.5, 14.5, 22. , 22. , 24. , 22.5, 29. , 24.5, 29. ,
33. , 20. , 18. , 18.5, 17.5, 29.5, 32. , 28. , 26.5, 20. , 13. ,
19. , 19. , 16.5, 16.5, 13. , 13. , 13. , 31.5, 30. , 36. , 25.5,
33.5, 17.5, 17. , 15.5, 15. , 17.5, 20.5, 19. , 18.5, 16. , 15.5,
15.5, 16. , 29. , 24.5, 26. , 25.5, 30.5, 33.5, 30. , 30.5, 22. ,
21.5, 21.5, 43.1, 36.1, 32.8, 39.4, 36.1, 19.9, 19.4, 20.2, 19.2,
20.5, 20.2, 25.1, 20.5, 19.4, 20.6, 20.8, 18.6, 18.1, 19.2, 17.7,
18.1, 17.5, 30. , 27.5, 27.2, 30.9, 21.1, 23.2, 23.8, 23.9, 20.3,
17. , 21.6, 16.2, 31.5, 29.5, 21.5, 19.8, 22.3, 20.2, 20.6, 17. ,
17.6, 16.5, 18.2, 16.9, 15.5, 19.2, 18.5, 31.9, 34.1, 35.7, 27.4,
25.4, 23. , 27.2, 23.9, 34.2, 34.5, 31.8, 37.3, 28.4, 28.8, 26.8,
33.5, 41.5, 38.1, 32.1, 37.2, 28. , 26.4, 24.3, 19.1, 34.3, 29.8,
31.3, 37. , 32.2, 46.6, 27.9, 40.8, 44.3, 43.4, 36.4, 30. , 44.6,
40.9, 33.8, 29.8, 32.7, 23.7, 35. , 23.6, 32.4, 27.2, 26.6, 25.8,
23.5, 30. , 39.1, 39. , 35.1, 32.3, 37. , 37.7, 34.1, 34.7, 34.4,
29.9, 33. , 34.5, 33.7, 32.4, 32.9, 31.6, 28.1, 30.7, 25.4, 24.2,
22.4, 26.6, 20.2, 17.6, 28. , 27. , 34. , 31. , 29. , 27. , 24. ,
23. , 36. , 37. , 31. , 38. , 36. , 36. , 36. , 34. , 38. , 32. ,
38. , 25. , 38. , 26. , 22. , 32. , 36. , 27. , 27. , 44. , 32. ,
28. , 31. ])
```

```

[118]: print("Mean of mpg = ", np.mean(mpg_data))
print("Median of mpg = ", np.median(mpg_data))
print("Standard deviation of mpg = ", np.std(mpg_data))
print("Number of cars with mpg>25 = ", np.count_nonzero(mpg_data>25))
```

```

Mean of mpg = 23.514572864321607
Median of mpg = 23.0
Standard deviation of mpg = 7.806159061274433
Number of cars with mpg>25 = 158
```

1.2 2. Filtering

Using NumPy, filter all cars with more than 6 cylinders.
Return the corresponding car name as a list.

```
[119]: cylinder_data = np.array([data[i][1] for i in range(1,len(data))]).astype('int')
      car_name_data = np.array([data[i][8] for i in range(1,len(data))])
```

```
[120]: car_with_more_than_6_cylinders = [car_name_data[i] for i in
      ↪range(len(car_name_data)) if cylinder_data[i]>6 ]
      print("Cars with with more than 6 cylinders = ",
      ↪len(car_with_more_than_6_cylinders))
```

Cars with with more than 6 cylinders = 103

```
[121]: car_with_more_than_6_cylinders
```

```
[121]: ['chevrolet chevelle malibu',
      'buick skylark 320',
      'plymouth satellite',
      'amc rebel sst',
      'ford torino',
      'ford galaxie 500',
      'chevrolet impala',
      'plymouth fury iii',
      'pontiac catalina',
      'amc ambassador dpl',
      'dodge challenger se',
      'plymouth 'cuda 340'',
      'chevrolet monte carlo',
      'buick estate wagon (sw)',
      'ford f250',
      'chevy c20',
      'dodge d200',
      'hi 1200d',
      'chevrolet impala',
      'pontiac catalina brougham',
      'ford galaxie 500',
      'plymouth fury iii',
      'dodge monaco (sw)',
      'ford country squire (sw)',
      'pontiac safari (sw)',
      'chevrolet impala',
      'pontiac catalina',
      'plymouth fury iii',
      'ford galaxie 500',
      'amc ambassador sst',
      'mercury marquis',
      'buick lesabre custom',
      'oldsmobile delta 88 royale',
      'chrysler newport royal',
      'amc matador (sw)',
      'chevrolet chevelle concours (sw)',
```

'ford gran torino (sw)',
'plymouth satellite custom (sw)',
'buick century 350',
'amc matador',
'chevrolet malibu',
'ford gran torino',
'dodge coronet custom',
'mercury marquis brougham',
'chevrolet caprice classic',
'ford ltd',
'plymouth fury gran sedan',
'chrysler new yorker brougham',
'buick electra 225 custom',
'amc ambassador brougham',
'chevrolet impala',
'ford country',
'plymouth custom suburb',
'oldsmobile vista cruiser',
'chevrolet monte carlo s',
'pontiac grand prix',
'dodge dart custom',
'oldsmobile omega',
'ford gran torino',
'buick century luxus (sw)',
'dodge coronet custom (sw)',
'ford gran torino (sw)',
'amc matador (sw)',
'pontiac catalina',
'chevrolet bel air',
'plymouth grand fury',
'ford ltd',
'chevrolet monza 2+2',
'ford mustang ii',
'chevrolet chevelle malibu classic',
'dodge coronet brougham',
'amc matador',
'ford gran torino',
'plymouth volare premier v8',
'cadillac seville',
'chevy c10',
'ford f108',
'dodge d100',
'chevrolet caprice classic',
'oldsmobile cutlass supreme',
'dodge monaco brougham',
'mercury cougar brougham',
'pontiac grand prix lj',

```
'chevrolet monte carlo landau',
'chrysler cordoba',
'ford thunderbird',
'oldsmobile cutlass salon brougham',
'dodge diplomat',
'mercury monarch ghia',
'chevrolet monte carlo landau',
'ford futura',
'dodge magnum xe',
'chevrolet caprice classic',
'ford ltd landau',
'mercury grand marquis',
'dodge st. regis',
'buick estate wagon (sw)',
'ford country squire (sw)',
'chevrolet malibu classic (sw)',
'chrysler lebaron town @ country (sw)',
'cadillac eldorado',
'oldsmobile cutlass salon brougham',
'oldsmobile cutlass ls']
```

1.3 3. Statistical Analysis

Compute the 25th, 50th, and 75th percentiles of the weight column using NumPy.

```
[122]: weight_data = np.array([data[i][4] for i in range(1,len(data))]).astype('int')
```

```
[123]: print("25th percentile of weight column = ", np.percentile(weight_data,25))
print("50th percentile of weight column = ", np.percentile(weight_data,50))
print("75th percentile of weight column = ", np.percentile(weight_data,75))
```

```
25th percentile of weight column = 2223.75
50th percentile of weight column = 2803.5
75th percentile of weight column = 3608.0
```

1.4 4. Array Manipulation

Convert the acceleration column into a NumPy array and normalize its values (scale between 0 and 1).

```
[124]: acceleration_data = np.array([data[i][5] for i in range(1,len(data))]).
      ↪astype('float64')
acceleration_data
```

```
[124]: array([12. , 11.5, 11. , 12. , 10.5, 10. ,  9. ,  8.5, 10. ,  8.5, 10. ,
           8. ,  9.5, 10. , 15. , 15.5, 15.5, 16. , 14.5, 20.5, 17.5, 14.5,
          17.5, 12.5, 15. , 14. , 15. , 13.5, 18.5, 14.5, 15.5, 14. , 19. ,
          13. , 15.5, 15.5, 15.5, 15.5, 12. , 11.5, 13.5, 13. , 11.5, 12. ,
```

```

12. , 13.5, 19. , 15. , 14.5, 14. , 14. , 19.5, 14.5, 19. , 18. ,
19. , 20.5, 15.5, 17. , 23.5, 19.5, 16.5, 12. , 12. , 13.5, 13. ,
11.5, 11. , 13.5, 13.5, 12.5, 13.5, 12.5, 14. , 16. , 14. , 14.5,
18. , 19.5, 18. , 16. , 17. , 14.5, 15. , 16.5, 13. , 11.5, 13. ,
14.5, 12.5, 11.5, 12. , 13. , 14.5, 11. , 11. , 11. , 16.5, 18. ,
16. , 16.5, 16. , 21. , 14. , 12.5, 13. , 12.5, 15. , 19. , 19.5,
16.5, 13.5, 18.5, 14. , 15.5, 13. , 9.5, 19.5, 15.5, 14. , 15.5,
11. , 14. , 13.5, 11. , 16.5, 17. , 16. , 17. , 19. , 16.5, 21. ,
17. , 17. , 18. , 16.5, 14. , 14.5, 13.5, 16. , 15.5, 16.5, 15.5,
14.5, 16.5, 19. , 14.5, 15.5, 14. , 15. , 15.5, 16. , 16. , 16. ,
21. , 19.5, 11.5, 14. , 14.5, 13.5, 21. , 18.5, 19. , 19. , 15. ,
13.5, 12. , 16. , 17. , 16. , 18.5, 13.5, 16.5, 17. , 14.5, 14. ,
17. , 15. , 17. , 14.5, 13.5, 17.5, 15.5, 16.9, 14.9, 17.7, 15.3,
13. , 13. , 13.9, 12.8, 15.4, 14.5, 17.6, 17.6, 22.2, 22.1, 14.2,
17.4, 17.7, 21. , 16.2, 17.8, 12.2, 17. , 16.4, 13.6, 15.7, 13.2,
21.9, 15.5, 16.7, 12.1, 12. , 15. , 14. , 18.5, 14.8, 18.6, 15.5,
16.8, 12.5, 19. , 13.7, 14.9, 16.4, 16.9, 17.7, 19. , 11.1, 11.4,
12.2, 14.5, 14.5, 16. , 18.2, 15.8, 17. , 15.9, 16.4, 14.1, 14.5,
12.8, 13.5, 21.5, 14.4, 19.4, 18.6, 16.4, 15.5, 13.2, 12.8, 19.2,
18.2, 15.8, 15.4, 17.2, 17.2, 15.8, 16.7, 18.7, 15.1, 13.2, 13.4,
11.2, 13.7, 16.5, 14.2, 14.7, 14.5, 14.8, 16.7, 17.6, 14.9, 15.9,
13.6, 15.7, 15.8, 14.9, 16.6, 15.4, 18.2, 17.3, 18.2, 16.6, 15.4,
13.4, 13.2, 15.2, 14.9, 14.3, 15. , 13. , 14. , 15.2, 14.4, 15. ,
20.1, 17.4, 24.8, 22.2, 13.2, 14.9, 19.2, 14.7, 16. , 11.3, 12.9,
13.2, 14.7, 18.8, 15.5, 16.4, 16.5, 18.1, 20.1, 18.7, 15.8, 15.5,
17.5, 15. , 15.2, 17.9, 14.4, 19.2, 21.7, 23.7, 19.9, 21.8, 13.8,
17.3, 18. , 15.3, 11.4, 12.5, 15.1, 14.3, 17. , 15.7, 16.4, 14.4,
12.6, 12.9, 16.9, 16.4, 16.1, 17.8, 19.4, 17.3, 16. , 14.9, 16.2,
20.7, 14.2, 15.8, 14.4, 16.8, 14.8, 18.3, 20.4, 19.6, 12.6, 13.8,
15.8, 19. , 17.1, 16.6, 19.6, 18.6, 18. , 16.2, 16. , 18. , 16.4,
20.5, 15.3, 18.2, 17.6, 14.7, 17.3, 14.5, 14.5, 16.9, 15. , 15.7,
16.2, 16.4, 17. , 14.5, 14.7, 13.9, 13. , 17.3, 15.6, 24.6, 11.6,
18.6, 19.4])

```

```

[125]: normalised_acceleration_data = acceleration_data/norm(acceleration_data)
print("Normalised acceleration values = \n", normalised_acceleration_data)

```

```

Normalised acceleration values =
[0.03804627 0.03646101 0.03487575 0.03804627 0.03329049 0.03170523
0.0285347 0.02694944 0.03170523 0.02694944 0.03170523 0.02536418
0.03011997 0.03170523 0.04755784 0.0491431 0.0491431 0.05072836
0.04597258 0.06499572 0.05548415 0.04597258 0.05548415 0.03963153
0.04755784 0.04438732 0.04755784 0.04280206 0.05865467 0.04597258
0.0491431 0.04438732 0.06023993 0.0412168 0.0491431 0.0491431
0.0491431 0.0491431 0.03804627 0.03646101 0.04280206 0.0412168
0.03646101 0.03804627 0.03804627 0.04280206 0.06023993 0.04755784
0.04597258 0.04438732 0.04438732 0.06182519 0.04597258 0.06023993
0.05706941 0.06023993 0.06499572 0.0491431 0.05389889 0.07450728]

```

0.06182519	0.05231362	0.03804627	0.03804627	0.04280206	0.0412168
0.03646101	0.03487575	0.04280206	0.04280206	0.03963153	0.04280206
0.03963153	0.04438732	0.05072836	0.04438732	0.04597258	0.05706941
0.06182519	0.05706941	0.05072836	0.05389889	0.04597258	0.04755784
0.05231362	0.0412168	0.03646101	0.0412168	0.04597258	0.03963153
0.03646101	0.03804627	0.0412168	0.04597258	0.03487575	0.03487575
0.03487575	0.05231362	0.05706941	0.05072836	0.05231362	0.05072836
0.06658098	0.04438732	0.03963153	0.0412168	0.03963153	0.04755784
0.06023993	0.06182519	0.05231362	0.04280206	0.05865467	0.04438732
0.0491431	0.0412168	0.03011997	0.06182519	0.0491431	0.04438732
0.0491431	0.03487575	0.04438732	0.04280206	0.03487575	0.05231362
0.05389889	0.05072836	0.05389889	0.06023993	0.05231362	0.06658098
0.05389889	0.05389889	0.05706941	0.05231362	0.04438732	0.04597258
0.04280206	0.05072836	0.0491431	0.05231362	0.0491431	0.04597258
0.05231362	0.06023993	0.04597258	0.0491431	0.04438732	0.04755784
0.0491431	0.05072836	0.05072836	0.05072836	0.06658098	0.06182519
0.03646101	0.04438732	0.04597258	0.04280206	0.06658098	0.05865467
0.06023993	0.06023993	0.04755784	0.04280206	0.03804627	0.05072836
0.05389889	0.05072836	0.05865467	0.04280206	0.05231362	0.05389889
0.04597258	0.04438732	0.05389889	0.04755784	0.05389889	0.04597258
0.04280206	0.05548415	0.0491431	0.05358183	0.04724079	0.05611825
0.048509	0.0412168	0.0412168	0.04407027	0.04058269	0.04882605
0.04597258	0.0558012	0.0558012	0.0703856	0.07006855	0.04502142
0.0551671	0.05611825	0.06658098	0.05136247	0.0564353	0.03868038
0.05389889	0.05199657	0.04311911	0.04977721	0.0418509	0.06943445
0.0491431	0.05294773	0.03836332	0.03804627	0.04755784	0.04438732
0.05865467	0.04692374	0.05897172	0.0491431	0.05326478	0.03963153
0.06023993	0.04343616	0.04724079	0.05199657	0.05358183	0.05611825
0.06023993	0.0351928	0.03614396	0.03868038	0.04597258	0.04597258
0.05072836	0.05770351	0.05009426	0.05389889	0.05041131	0.05199657
0.04470437	0.04597258	0.04058269	0.04280206	0.06816624	0.04565553
0.06150814	0.05897172	0.05199657	0.0491431	0.0418509	0.04058269
0.06087404	0.05770351	0.05009426	0.04882605	0.05453299	0.05453299
0.05009426	0.05294773	0.05928877	0.04787489	0.0418509	0.042485
0.03550985	0.04343616	0.05231362	0.04502142	0.04660668	0.04597258
0.04692374	0.05294773	0.0558012	0.04724079	0.05041131	0.04311911
0.04977721	0.05009426	0.04724079	0.05263068	0.04882605	0.05770351
0.05485004	0.05770351	0.05263068	0.04882605	0.042485	0.0418509
0.04819195	0.04724079	0.04533847	0.04755784	0.0412168	0.04438732
0.04819195	0.04565553	0.04755784	0.06372751	0.0551671	0.07862896
0.0703856	0.0418509	0.04724079	0.06087404	0.04660668	0.05072836
0.03582691	0.04089974	0.0418509	0.04660668	0.05960583	0.0491431
0.05199657	0.05231362	0.05738646	0.06372751	0.05928877	0.05009426
0.0491431	0.05548415	0.04755784	0.04819195	0.05675236	0.04565553
0.06087404	0.06880034	0.07514139	0.0630934	0.06911739	0.04375321
0.05485004	0.05706941	0.048509	0.03614396	0.03963153	0.04787489
0.04533847	0.05389889	0.04977721	0.05199657	0.04565553	0.03994859
0.04089974	0.05358183	0.05199657	0.05104542	0.0564353	0.06150814

```

0.05485004 0.05072836 0.04724079 0.05136247 0.06562982 0.04502142
0.05009426 0.04565553 0.05326478 0.04692374 0.05802057 0.06467866
0.06214224 0.03994859 0.04375321 0.05009426 0.06023993 0.05421594
0.05263068 0.06214224 0.05897172 0.05706941 0.05136247 0.05072836
0.05706941 0.05199657 0.06499572 0.048509 0.05770351 0.0558012
0.04660668 0.05485004 0.04597258 0.04597258 0.05358183 0.04755784
0.04977721 0.05136247 0.05199657 0.05389889 0.04597258 0.04660668
0.04407027 0.0412168 0.05485004 0.04946015 0.07799486 0.03677806
0.05897172 0.06150814]

```

```

[126]: normalised_acceleration_data_min_max = (acceleration_data-np.
        ↪min(acceleration_data))/(np.max(acceleration_data)-np.min(acceleration_data))
print("Normalised acceleration values using min-max scaling = \n",
        ↪normalised_acceleration_data_min_max)

```

```

Normalised acceleration values using min-max scaling =
[0.23809524 0.20833333 0.17857143 0.23809524 0.14880952 0.11904762
0.05952381 0.0297619 0.11904762 0.0297619 0.11904762 0.
0.08928571 0.11904762 0.41666667 0.44642857 0.44642857 0.47619048
0.38690476 0.74404762 0.56547619 0.38690476 0.56547619 0.26785714
0.41666667 0.35714286 0.41666667 0.32738095 0.625 0.38690476
0.44642857 0.35714286 0.6547619 0.29761905 0.44642857 0.44642857
0.44642857 0.44642857 0.23809524 0.20833333 0.32738095 0.29761905
0.20833333 0.23809524 0.23809524 0.32738095 0.6547619 0.41666667
0.38690476 0.35714286 0.35714286 0.68452381 0.38690476 0.6547619
0.5952381 0.6547619 0.74404762 0.44642857 0.53571429 0.92261905
0.68452381 0.50595238 0.23809524 0.23809524 0.32738095 0.29761905
0.20833333 0.17857143 0.32738095 0.32738095 0.26785714 0.32738095
0.26785714 0.35714286 0.47619048 0.35714286 0.38690476 0.5952381
0.68452381 0.5952381 0.47619048 0.53571429 0.38690476 0.41666667
0.50595238 0.29761905 0.20833333 0.29761905 0.38690476 0.26785714
0.20833333 0.23809524 0.29761905 0.38690476 0.17857143 0.17857143
0.17857143 0.50595238 0.5952381 0.47619048 0.50595238 0.47619048
0.77380952 0.35714286 0.26785714 0.29761905 0.26785714 0.41666667
0.6547619 0.68452381 0.50595238 0.32738095 0.625 0.35714286
0.44642857 0.29761905 0.08928571 0.68452381 0.44642857 0.35714286
0.44642857 0.17857143 0.35714286 0.32738095 0.17857143 0.50595238
0.53571429 0.47619048 0.53571429 0.6547619 0.50595238 0.77380952
0.53571429 0.53571429 0.5952381 0.50595238 0.35714286 0.38690476
0.32738095 0.47619048 0.44642857 0.50595238 0.44642857 0.38690476
0.50595238 0.6547619 0.38690476 0.44642857 0.35714286 0.41666667
0.44642857 0.47619048 0.47619048 0.47619048 0.77380952 0.68452381
0.20833333 0.35714286 0.38690476 0.32738095 0.77380952 0.625
0.6547619 0.6547619 0.41666667 0.32738095 0.23809524 0.47619048
0.53571429 0.47619048 0.625 0.32738095 0.50595238 0.53571429
0.38690476 0.35714286 0.53571429 0.41666667 0.53571429 0.38690476
0.32738095 0.56547619 0.44642857 0.5297619 0.41071429 0.57738095
0.43452381 0.29761905 0.29761905 0.35119048 0.28571429 0.44047619

```



```

0.38690476 0.57142857 0.57142857 0.8452381 0.83928571 0.36904762
0.55952381 0.57738095 0.77380952 0.48809524 0.58333333 0.25
0.53571429 0.5 0.33333333 0.45833333 0.30952381 0.82738095
0.44642857 0.51785714 0.24404762 0.23809524 0.41666667 0.35714286
0.625 0.4047619 0.63095238 0.44642857 0.52380952 0.26785714
0.6547619 0.33928571 0.41071429 0.5 0.5297619 0.57738095
0.6547619 0.18452381 0.20238095 0.25 0.38690476 0.38690476
0.47619048 0.60714286 0.46428571 0.53571429 0.4702381 0.5
0.36309524 0.38690476 0.28571429 0.32738095 0.80357143 0.38095238
0.67857143 0.63095238 0.5 0.44642857 0.30952381 0.28571429
0.66666667 0.60714286 0.46428571 0.44047619 0.54761905 0.54761905
0.46428571 0.51785714 0.63690476 0.42261905 0.30952381 0.32142857
0.19047619 0.33928571 0.50595238 0.36904762 0.39880952 0.38690476
0.4047619 0.51785714 0.57142857 0.41071429 0.4702381 0.33333333
0.45833333 0.46428571 0.41071429 0.51190476 0.44047619 0.60714286
0.55357143 0.60714286 0.51190476 0.44047619 0.32142857 0.30952381
0.42857143 0.41071429 0.375 0.41666667 0.29761905 0.35714286
0.42857143 0.38095238 0.41666667 0.7202381 0.55952381 1.
0.8452381 0.30952381 0.41071429 0.66666667 0.39880952 0.47619048
0.19642857 0.29166667 0.30952381 0.39880952 0.64285714 0.44642857
0.5 0.50595238 0.60119048 0.7202381 0.63690476 0.46428571
0.44642857 0.56547619 0.41666667 0.42857143 0.58928571 0.38095238
0.66666667 0.81547619 0.93452381 0.70833333 0.82142857 0.3452381
0.55357143 0.5952381 0.43452381 0.20238095 0.26785714 0.42261905
0.375 0.53571429 0.45833333 0.5 0.38095238 0.27380952
0.29166667 0.5297619 0.5 0.48214286 0.58333333 0.67857143
0.55357143 0.47619048 0.41071429 0.48809524 0.75595238 0.36904762
0.46428571 0.38095238 0.52380952 0.4047619 0.61309524 0.73809524
0.69047619 0.27380952 0.3452381 0.46428571 0.6547619 0.54166667
0.51190476 0.69047619 0.63095238 0.5952381 0.48809524 0.47619048
0.5952381 0.5 0.74404762 0.43452381 0.60714286 0.57142857
0.39880952 0.55357143 0.38690476 0.38690476 0.5297619 0.41666667
0.45833333 0.48809524 0.5 0.53571429 0.38690476 0.39880952
0.35119048 0.29761905 0.55357143 0.45238095 0.98809524 0.21428571
0.63095238 0.67857143]

```

1.5 5. Broadcasting

Increase all horsepower values by 10% and store the updated values in a new NumPy array. Handle missing data (if any) by replacing it with the mean of the column before applying the increase.

```
[127]: horse_power_data = np.array([data[i][3] for i in range(1,len(data))])
horse_power_data
```

```
[127]: array(['130', '165', '150', '150', '140', '198', '220', '215', '225',
            '190', '170', '160', '150', '225', '95', '95', '97', '85', '88',
            '46', '87', '90', '95', '113', '90', '215', '200', '210', '193',
            '88', '90', '95', '?', '100', '105', '100', '88', '100', '165',
```

```

'175', '153', '150', '180', '170', '175', '110', '72', '100', '88',
'86', '90', '70', '76', '65', '69', '60', '70', '95', '80', '54',
'90', '86', '165', '175', '150', '153', '150', '208', '155', '160',
'190', '97', '150', '130', '140', '150', '112', '76', '87', '69',
'86', '92', '97', '80', '88', '175', '150', '145', '137', '150',
'198', '150', '158', '150', '215', '225', '175', '105', '100',
'100', '88', '95', '46', '150', '167', '170', '180', '100', '88',
'72', '94', '90', '85', '107', '90', '145', '230', '49', '75',
'91', '112', '150', '110', '122', '180', '95', '?', '100', '100',
'67', '80', '65', '75', '100', '110', '105', '140', '150', '150',
'140', '150', '83', '67', '78', '52', '61', '75', '75', '75', '97',
'93', '67', '95', '105', '72', '72', '170', '145', '150', '148',
'110', '105', '110', '95', '110', '110', '129', '75', '83', '100',
'78', '96', '71', '97', '97', '70', '90', '95', '88', '98', '115',
'53', '86', '81', '92', '79', '83', '140', '150', '120', '152',
'100', '105', '81', '90', '52', '60', '70', '53', '100', '78',
'110', '95', '71', '70', '75', '72', '102', '150', '88', '108',
'120', '180', '145', '130', '150', '68', '80', '58', '96', '70',
'145', '110', '145', '130', '110', '105', '100', '98', '180',
'170', '190', '149', '78', '88', '75', '89', '63', '83', '67',
'78', '97', '110', '110', '48', '66', '52', '70', '60', '110',
'140', '139', '105', '95', '85', '88', '100', '90', '105', '85',
'110', '120', '145', '165', '139', '140', '68', '95', '97', '75',
'95', '105', '85', '97', '103', '125', '115', '133', '71', '68',
'115', '85', '88', '90', '110', '130', '129', '138', '135', '155',
'142', '125', '150', '71', '65', '80', '80', '77', '125', '71',
'90', '70', '70', '65', '69', '90', '115', '115', '90', '76', '60',
'70', '65', '90', '88', '90', '90', '78', '90', '75', '92', '75',
'65', '105', '65', '48', '48', '67', '67', '67', '?', '67', '62',
'132', '100', '88', '?', '72', '84', '84', '92', '110', '84', '58',
'64', '60', '67', '65', '62', '68', '63', '65', '65', '74', '?',
'75', '75', '100', '74', '80', '76', '116', '120', '110', '105',
'88', '85', '88', '88', '88', '85', '84', '90', '92', '?', '74',
'68', '68', '63', '70', '88', '75', '70', '67', '67', '67', '110',
'85', '92', '112', '96', '84', '90', '86', '52', '84', '79', '82'],
dtype='<U3')

```

```

[128]: mean_sum = 0
count = 0
for i in range(len(horse_power_data)):
    if(horse_power_data[i] != '?'):
        count += 1
        mean_sum += horse_power_data[i].astype('int')
mean = mean_sum/count

for i in range(len(horse_power_data)):
    if(horse_power_data[i] == '?'):

```

```

horse_power_data[i] = np.round(mean)

horse_power_data = horse_power_data.astype('int')

horse_power_data

```

```

[128]: array([130, 165, 150, 150, 140, 198, 220, 215, 225, 190, 170, 160, 150,
225, 95, 95, 97, 85, 88, 46, 87, 90, 95, 113, 90, 215,
200, 210, 193, 88, 90, 95, 104, 100, 105, 100, 88, 100, 165,
175, 153, 150, 180, 170, 175, 110, 72, 100, 88, 86, 90, 70,
76, 65, 69, 60, 70, 95, 80, 54, 90, 86, 165, 175, 150,
153, 150, 208, 155, 160, 190, 97, 150, 130, 140, 150, 112, 76,
87, 69, 86, 92, 97, 80, 88, 175, 150, 145, 137, 150, 198,
150, 158, 150, 215, 225, 175, 105, 100, 100, 88, 95, 46, 150,
167, 170, 180, 100, 88, 72, 94, 90, 85, 107, 90, 145, 230,
49, 75, 91, 112, 150, 110, 122, 180, 95, 104, 100, 100, 67,
80, 65, 75, 100, 110, 105, 140, 150, 150, 140, 150, 83, 67,
78, 52, 61, 75, 75, 75, 97, 93, 67, 95, 105, 72, 72,
170, 145, 150, 148, 110, 105, 110, 95, 110, 110, 129, 75, 83,
100, 78, 96, 71, 97, 97, 70, 90, 95, 88, 98, 115, 53,
86, 81, 92, 79, 83, 140, 150, 120, 152, 100, 105, 81, 90,
52, 60, 70, 53, 100, 78, 110, 95, 71, 70, 75, 72, 102,
150, 88, 108, 120, 180, 145, 130, 150, 68, 80, 58, 96, 70,
145, 110, 145, 130, 110, 105, 100, 98, 180, 170, 190, 149, 78,
88, 75, 89, 63, 83, 67, 78, 97, 110, 110, 48, 66, 52,
70, 60, 110, 140, 139, 105, 95, 85, 88, 100, 90, 105, 85,
110, 120, 145, 165, 139, 140, 68, 95, 97, 75, 95, 105, 85,
97, 103, 125, 115, 133, 71, 68, 115, 85, 88, 90, 110, 130,
129, 138, 135, 155, 142, 125, 150, 71, 65, 80, 80, 77, 125,
71, 90, 70, 70, 65, 69, 90, 115, 115, 90, 76, 60, 70,
65, 90, 88, 90, 90, 78, 90, 75, 92, 75, 65, 105, 65,
48, 48, 67, 67, 67, 104, 67, 62, 132, 100, 88, 104, 72,
84, 84, 92, 110, 84, 58, 64, 60, 67, 65, 62, 68, 63,
65, 65, 74, 104, 75, 75, 100, 74, 80, 76, 116, 120, 110,
105, 88, 85, 88, 88, 88, 85, 84, 90, 92, 104, 74, 68,
68, 63, 70, 88, 75, 70, 67, 67, 67, 110, 85, 92, 112,
96, 84, 90, 86, 52, 84, 79, 82])

```

```

[129]: increased_horse_power_data = horse_power_data+horse_power_data*0.1
print("Increased horsepower (+10%) = \n", increased_horse_power_data)

```

```

Increased horsepower (+10%) =
[143.  181.5 165.  165.  154.  217.8 242.  236.5 247.5 209.  187.  176.
165.  247.5 104.5 104.5 106.7 93.5  96.8  50.6  95.7  99.  104.5 124.3
99.  236.5 220.  231.  212.3 96.8  99.  104.5 114.4 110.  115.5 110.
96.8 110.  181.5 192.5 168.3 165.  198.  187.  192.5 121.  79.2 110.
96.8  94.6  99.   77.   83.6  71.5  75.9  66.   77.  104.5  88.   59.4
99.   94.6 181.5 192.5 165.  168.3 165.  228.8 170.5 176.  209.  106.7

```

```

165.  143.  154.  165.  123.2  83.6  95.7  75.9  94.6  101.2  106.7  88.
 96.8 192.5 165.  159.5 150.7 165.  217.8 165.  173.8 165.  236.5 247.5
192.5 115.5 110.  110.   96.8 104.5  50.6 165.  183.7 187.  198.  110.
 96.8  79.2 103.4  99.   93.5 117.7  99.  159.5 253.   53.9  82.5 100.1
123.2 165.  121.  134.2 198.  104.5 114.4 110.  110.   73.7  88.   71.5
 82.5 110.  121.  115.5 154.  165.  165.  154.  165.   91.3  73.7  85.8
 57.2  67.1  82.5  82.5  82.5 106.7 102.3  73.7 104.5 115.5  79.2  79.2
187.  159.5 165.  162.8 121.  115.5 121.  104.5 121.  121.  141.9  82.5
 91.3 110.   85.8 105.6  78.1 106.7 106.7  77.   99.  104.5  96.8 107.8
126.5  58.3  94.6  89.1 101.2  86.9  91.3 154.  165.  132.  167.2 110.
115.5  89.1  99.   57.2  66.   77.   58.3 110.   85.8 121.  104.5  78.1
 77.   82.5  79.2 112.2 165.   96.8 118.8 132.  198.  159.5 143.  165.
 74.8  88.   63.8 105.6  77.  159.5 121.  159.5 143.  121.  115.5 110.
107.8 198.  187.  209.  163.9  85.8  96.8  82.5  97.9  69.3  91.3  73.7
 85.8 106.7 121.  121.   52.8  72.6  57.2  77.   66.  121.  154.  152.9
115.5 104.5  93.5  96.8 110.   99.  115.5  93.5 121.  132.  159.5 181.5
152.9 154.   74.8 104.5 106.7  82.5 104.5 115.5  93.5 106.7 113.3 137.5
126.5 146.3  78.1  74.8 126.5  93.5  96.8  99.  121.  143.  141.9 151.8
148.5 170.5 156.2 137.5 165.   78.1  71.5  88.   88.   84.7 137.5  78.1
 99.   77.   77.   71.5  75.9  99.  126.5 126.5  99.   83.6  66.   77.
 71.5  99.   96.8  99.   99.   85.8  99.   82.5 101.2  82.5  71.5 115.5
 71.5  52.8  52.8  73.7  73.7  73.7 114.4  73.7  68.2 145.2 110.   96.8
114.4  79.2  92.4  92.4 101.2 121.   92.4  63.8  70.4  66.   73.7  71.5
 68.2  74.8  69.3  71.5  71.5  81.4 114.4  82.5  82.5 110.   81.4  88.
 83.6 127.6 132.  121.  115.5  96.8  93.5  96.8  96.8  96.8  93.5  92.4
 99.  101.2 114.4  81.4  74.8  74.8  69.3  77.   96.8  82.5  77.   73.7
 73.7  73.7 121.   93.5 101.2 123.2 105.6  92.4  99.   94.6  57.2  92.4
 86.9  90.2]

```

1.6 6. Boolean Indexing

Find the average displacement of cars with an origin of 2 (Europe) using NumPy indexing.

```
[130]: displacement_data = np.array([data[i][2] for i in range(1,len(data))])
      origin_data = np.array([data[i][7] for i in range(1,len(data))]).astype('int')
```

```
[131]: origin_boolean_mask = origin_data == 2
```

```
[132]: print("Average displacement of cars with origin 2 (Europe) = ", np.round(np.
      ↪mean(displacement_data[origin_boolean_mask].astype('float64')),2))
```

Average displacement of cars with origin 2 (Europe) = 109.14

1.7 7. Matrix Operations

Create a 2D NumPy array containing the columns mpg, horsepower, and weight.
Compute the dot product of this matrix with a given vector [1, 0.5, -0.2].

```
[133]: d2_array = np.array([mpg_data, horse_power_data, weight_data])
print(d2_array)
print(d2_array.shape)
```

```
[[ 18.   15.   18. ...   32.   28.   31.]
 [130. 165. 150. ...   84.   79.   82.]
 [3504. 3693. 3436. ... 2295. 2625. 2720.]]
(3, 398)
```

```
[134]: vector = np.array([1, 0.5, -0.2])
print(vector)
print(vector.shape)
```

```
[ 1.    0.5 -0.2]
(3,)
```

```
[135]: dot_product = np.dot(vector, d2_array)
print("Dot product = ", dot_product)
print(dot_product.shape)
```

```
Dot product =  [-617.8 -641.1 -594.2 -595.6 -602.8 -754.2 -746.8 -740.9 -758.5
-660.
-612.6 -627.8 -662.2 -490.7 -402.9 -497.1 -488.3 -453.9 -355.  -318.
-465.9 -417.  -402.5 -364.3 -463.6 -805.5 -765.2 -760.4 -840.9 -355.
-379.8 -373.1 -332.2 -457.8 -619.3 -598.8 -597.4 -589.6 -745.3 -791.3
-740.3 -730.2 -889.  -851.2 -927.5 -519.4 -423.6 -587.4 -565.8 -378.
-351.6 -349.8 -345.  -291.1 -253.1 -309.8 -330.  -384.1 -360.2 -400.8
-416.6 -381.2 -759.3 -775.5 -737.  -735.3 -642.4 -811.6 -809.9 -799.2
-776.4 -398.5 -688.4 -741.6 -775.8 -726.4 -512.6 -442.2 -531.3 -377.3
-414.  -383.6 -429.7 -364.8 -349.  -719.5 -645.4 -712.1 -725.9 -665.4
-879.4 -804.8 -780.6 -758.4 -826.5 -865.7 -663.7 -553.7 -589.6 -521.
-542.2 -510.3 -341.  -913.4 -885.7 -832.8 -797.8 -489.8 -391.8 -423.2
-406.8 -361.8 -400.5 -419.9 -382.  -728.9 -724.6 -319.9 -370.1 -450.9
-498.6 -589.8 -453.  -480.4 -631.8 -552.9 -502.  -511.2 -602.2 -325.5
-424.2 -302.7 -445.9 -690.2 -655.4 -652.1 -742.2 -851.8 -802.4 -843.6
-762.4 -373.3 -333.1 -395.  -272.8 -338.1 -359.5 -360.1 -385.7 -425.3
-405.7 -335.5 -586.3 -621.3 -635.4 -580.6 -832.6 -800.5 -808.6 -843.4
-709.4 -710.9 -676.  -691.5 -531.8 -569.2 -556.3 -367.7 -463.3 -512.8
-456.4 -468.4 -384.1 -436.5 -530.3 -323.4 -578.2 -468.3 -524.4 -518.
-451.7 -299.5 -421.8 -378.5 -443.4 -385.5 -371.9 -755.5 -747.  -716.9
-752.5 -574.6 -596.1 -537.9 -549.5 -352.  -378.3 -323.4 -299.5 -660.2
-657.8 -655.5 -573.6 -300.  -331.  -365.5 -450.5 -559.  -700.  -591.
-513.  -687.5 -769.5 -725.5 -696.  -663.  -343.5 -361.  -300.  -386.5
-320.5 -686.  -740.  -740.  -779.  -631.5 -612.  -657.  -637.5 -738.
-732.5 -754.5 -776.5 -320.  -479.5 -389.5 -481.  -348.2 -340.  -333.5
-368.5 -492.5 -443.5 -467.5 -329.9 -290.9 -338.2 -339.6 -293.9 -598.1
-657.6 -624.3 -635.3 -563.  -530.3 -474.9 -615.5 -577.6 -602.9 -550.7
-650.4 -603.9 -593.3 -588.8 -553.4 -728.5 -367.  -437.  -384.3 -377.6
-434.4 -473.3 -504.7 -408.6 -494.2 -548.5 -479.9 -599.3 -331.  -363.5
```

```

-570.  -535.7 -511.7 -587.8 -596.4 -686.  -662.9 -705.5 -680.3 -777.6
-724.3 -639.3 -694.5 -317.6 -328.4 -307.3 -466.6 -642.1 -694.5 -575.3
-615.1 -370.8 -360.5 -339.7 -354.2 -460.6 -432.7 -455.7 -432.7 -349.3
-325.5 -356.9 -334.1 -462.6 -503.6 -531.3 -612.1 -364.3 -467.4 -439.6
-403.8 -383.3 -342.9 -479.6 -348.7 -348.7 -399.6 -520.1 -586.5 -291.9
-274.1 -361.7 -308.2 -483.3 -410.3 -421.  -505.4 -389.6 -428.8 -458.4
-452.2 -466.5 -405.  -282.9 -304.  -286.9 -347.2 -325.5 -341.3 -328.9
-376.8 -342.1 -413.6 -368.  -377.5 -370.8 -400.1 -440.1 -458.4 -577.9
-563.3 -496.6 -501.8 -605.6 -665.9 -547.8 -632.9 -449.  -457.  -401.
-441.5 -434.  -475.  -503.  -532.  -323.  -334.  -329.  -355.5 -354.
-352.  -367.5 -380.  -321.5 -327.5 -327.5 -509.  -522.5 -445.  -489.
-453.  -396.  -518.  -488.  -356.  -385.  -457.5 -472. ]
(398,)

```

1.8 8. Sorting

Use NumPy to sort the cars by model_year in descending order and display the first five car names.

```
[136]: model_year_data = np.array([data[i][6] for i in range(1,len(data))])
model_year_car_name_data = np.array([model_year_data, car_name_data])
```

```
[137]: desc_sorted_model_year_car_name_data = np.sort(np.
↳column_stack(([model_year_data, car_name_data])))[::-1]
```

```
[138]: for i in range(5):
print(desc_sorted_model_year_car_name_data[i][1])
```

```

chevy s-10
ford ranger
dodge rampage
vw pickup
ford mustang gl

```

1.9 9. Correlation

Compute the Pearson correlation coefficient between mpg and weight using NumPy.

```
[139]: print("Pearson Correlation Coefficient = \n", np.corrcoef(mpg_data,↳
↳weight_data))
```

```

Pearson Correlation Coefficient =
[[ 1.          -0.83174093]
 [-0.83174093  1.          ]]

```

1.10 10. Conditional Aggregates

Calculate the mean mpg for cars grouped by the number of cylinders using NumPy techniques.

```
[140]: mpg_cylinder_data = np.column_stack([mpg_data, cylinder_data])
```

```
[141]: unique_cylinders = np.unique(np.delete(np.column_stack((mpg_data,
↪cylinder_data))), 0, 1))
unique_cylinders
```

```
[141]: array([3., 4., 5., 6., 8.])
```

```
[142]: for i in unique_cylinders:
        sum_cars = 0
        sum_mpg = 0
        mean = 0
        for j in range(len(mpg_cylinder_data)):
            if (mpg_cylinder_data[j][1]==i):
                sum_cars+=1
                sum_mpg+=mpg_cylinder_data[j][0]
        mean = sum_mpg/sum_cars
        print("Mean mpg of",i,"cylinder cars = ", mean)
```

```
Mean mpg of 3.0 cylinder cars = 20.55
Mean mpg of 4.0 cylinder cars = 29.28676470588236
Mean mpg of 5.0 cylinder cars = 27.366666666666664
Mean mpg of 6.0 cylinder cars = 19.985714285714284
Mean mpg of 8.0 cylinder cars = 14.963106796116508
```

2 Pandas Problems

```
[143]: import pandas as pd
```

2.1 1. Basic Exploration

Load the dataset into a Pandas DataFrame. Display: - The first 10 rows - The total number of rows and columns - Summary statistics for numerical columns

```
[144]: df = pd.read_csv("auto-mpg.csv")
```

```
[145]: df.dtypes
```

```
[145]: mpg                float64
cylinders                int64
displacement            float64
horsepower              object
weight                  int64
acceleration            float64
model year              int64
origin                  int64
car name                object
dtype: object
```

```
[146]: df['horsepower'] = df['horsepower'].apply(pd.to_numeric, errors='coerce')
df.dtypes
```

```
[146]: mpg                float64
cylinders              int64
displacement          float64
horsepower            float64
weight               int64
acceleration         float64
model year           int64
origin              int64
car name             object
dtype: object
```

```
[147]: # First 10 rows
df.head(10)
```

```
[147]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	18.0	8	307.0	130.0	3504	12.0	
1	15.0	8	350.0	165.0	3693	11.5	
2	18.0	8	318.0	150.0	3436	11.0	
3	16.0	8	304.0	150.0	3433	12.0	
4	17.0	8	302.0	140.0	3449	10.5	
5	15.0	8	429.0	198.0	4341	10.0	
6	14.0	8	454.0	220.0	4354	9.0	
7	14.0	8	440.0	215.0	4312	8.5	
8	14.0	8	455.0	225.0	4425	10.0	
9	15.0	8	390.0	190.0	3850	8.5	

	model year	origin	car name
0	70	1	chevrolet chevelle malibu
1	70	1	buick skylark 320
2	70	1	plymouth satellite
3	70	1	amc rebel sst
4	70	1	ford torino
5	70	1	ford galaxie 500
6	70	1	chevrolet impala
7	70	1	plymouth fury iii
8	70	1	pontiac catalina
9	70	1	amc ambassador dpl

```
[148]: # Total number of rows and columns
print("No. of rows = ", df.shape[0])
print("No. of columns = ", df.shape[1])
```

```
No. of rows = 398
No. of columns = 9
```



```
[149]: # Summary statistics for numerical columns
df.describe()
```

```
[149]:
```

	mpg	cylinders	displacement	horsepower	weight \
count	398.000000	398.000000	398.000000	392.000000	398.000000
mean	23.514573	5.454774	193.425879	104.469388	2970.424623
std	7.815984	1.701004	104.269838	38.491160	846.841774
min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

	acceleration	model year	origin
count	398.000000	398.000000	398.000000
mean	15.568090	76.010050	1.572864
std	2.757689	3.697627	0.802055
min	8.000000	70.000000	1.000000
25%	13.825000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.175000	79.000000	2.000000
max	24.800000	82.000000	3.000000

2.2 2. Filtering and Indexing

Find all cars manufactured in 1975 with a weight less than 3000. Return the DataFrame with selected columns: car_name, weight, and mpg.

```
[150]: filtered_df = df[(df["model year"]==75) & (df["weight"]<3000)]
filtered_df[["car name","weight","mpg"]]
```

```
[150]:
```

	car name	weight	mpg
167	toyota corolla	2171	29.0
168	ford pinto	2639	23.0
169	amc gremlin	2914	20.0
170	pontiac astro	2592	23.0
171	toyota corona	2702	24.0
172	volkswagen dasher	2223	25.0
173	datsum 710	2545	24.0
174	ford pinto	2984	18.0
175	volkswagen rabbit	1937	29.0
177	audi 100ls	2694	23.0
178	peugeot 504	2957	23.0
179	volvo 244dl	2945	22.0
180	saab 99le	2671	25.0
181	honda civic cvcc	1795	33.0

2.3 3. Handling Missing Data

Identify if there are any missing values in the dataset. Replace missing values in the horsepower column with the column's median.

```
[151]: df.isna().any()
```

```
[151]: mpg           False
      cylinders      False
      displacement  False
      horsepower     True
      weight         False
      acceleration  False
      model year     False
      origin         False
      car name       False
      dtype: bool
```

```
[152]: df['horsepower'] = df['horsepower'].fillna(df['horsepower'].mean())
```

```
[153]: df.isna().any()
```

```
[153]: mpg           False
      cylinders      False
      displacement  False
      horsepower     False
      weight         False
      acceleration  False
      model year     False
      origin         False
      car name       False
      dtype: bool
```

2.4 4. Data Transformation

Add a new column power_to_weight_ratio, calculated as horsepower/weight.

```
[154]: df['power_to_weight_ratio'] = df['horsepower']/df['weight']
```

```
[155]: df
```

```
[155]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	18.0	8	307.0	130.0	3504	12.0	
1	15.0	8	350.0	165.0	3693	11.5	
2	18.0	8	318.0	150.0	3436	11.0	
3	16.0	8	304.0	150.0	3433	12.0	
4	17.0	8	302.0	140.0	3449	10.5	
..	
393	27.0	4	140.0	86.0	2790	15.6	

394	44.0	4	97.0	52.0	2130	24.6
395	32.0	4	135.0	84.0	2295	11.6
396	28.0	4	120.0	79.0	2625	18.6
397	31.0	4	119.0	82.0	2720	19.4

	model	year	origin	car name	power_to_weight_ratio
0		70	1	chevrolet chevelle malibu	0.037100
1		70	1	buick skylark 320	0.044679
2		70	1	plymouth satellite	0.043655
3		70	1	amc rebel sst	0.043694
4		70	1	ford torino	0.040591
..
393		82	1	ford mustang gl	0.030824
394		82	2	vw pickup	0.024413
395		82	1	dodge rampage	0.036601
396		82	1	ford ranger	0.030095
397		82	1	chevy s-10	0.030147

[398 rows x 10 columns]

2.5 5. Group By

Group the cars by origin and calculate the mean mpg for each group.

```
[156]: group = df.groupby(['origin'])
```

```
[157]: group['mpg'].mean()
```

```
[157]: origin
1    20.083534
2    27.891429
3    30.450633
Name: mpg, dtype: float64
```

2.6 6. Sorting

Sort the DataFrame by mpg in descending order and display the top 10 cars with the highest mpg.

```
[158]: df.sort_values(by='mpg', ascending=False).head(10)
```

```
[158]:      mpg  cylinders  displacement  horsepower  weight  acceleration  \
322  46.6         4         86.0     65.000000   2110         17.9
329  44.6         4         91.0     67.000000   1850         13.8
325  44.3         4         90.0     48.000000   2085         21.7
394  44.0         4         97.0     52.000000   2130         24.6
326  43.4         4         90.0     48.000000   2335         23.7
244  43.1         4         90.0     48.000000   1985         21.5
309  41.5         4         98.0     76.000000   2144         14.7
```

330	40.9	4	85.0	104.469388	1835	17.3
324	40.8	4	85.0	65.000000	2110	19.2
247	39.4	4	85.0	70.000000	2070	18.6

	model	year	origin	car name	\
322		80	3	mazda glc	
329		80	3	honda civic 1500 gl	
325		80	2	vw rabbit c (diesel)	
394		82	2	vw pickup	
326		80	2	vw dasher (diesel)	
244		78	2	volkswagen rabbit custom diesel	
309		80	2	vw rabbit	
330		80	2	renault lecar deluxe	
324		80	3	datsum 210	
247		78	3	datsum b210 gx	

	power_to_weight_ratio
322	0.030806
329	0.036216
325	0.023022
394	0.024413
326	0.020557
244	0.024181
309	0.035448
330	0.056932
324	0.030806
247	0.033816

2.7 7. Apply Function

Create a new column `performance_score` using a custom function:

```
def performance_score(row):
```

```
    return row['mpg'] * row['acceleration'] / row['weight']
```

Apply this function to each row and store the result in the new column.

```
[159]: def performance_score(row):
        return row['mpg']*row['acceleration']/row['weight']
```

```
[160]: df['performance_score']=performance_score(df)
```

```
[161]: df
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	18.0	8	307.0	130.0	3504	12.0	
1	15.0	8	350.0	165.0	3693	11.5	
2	18.0	8	318.0	150.0	3436	11.0	
3	16.0	8	304.0	150.0	3433	12.0	
4	17.0	8	302.0	140.0	3449	10.5	

```

..      ...      ...      ...      ...      ...      ...
393  27.0      4      140.0      86.0      2790      15.6
394  44.0      4      97.0      52.0      2130      24.6
395  32.0      4      135.0      84.0      2295      11.6
396  28.0      4      120.0      79.0      2625      18.6
397  31.0      4      119.0      82.0      2720      19.4

      model year  origin      car name  power_to_weight_ratio \
0          70      1  chevrolet chevelle malibu      0.037100
1          70      1      buick skylark 320      0.044679
2          70      1  plymouth satellite      0.043655
3          70      1      amc rebel sst      0.043694
4          70      1      ford torino      0.040591
..      ...      ...      ...      ...
393          82      1  ford mustang gl      0.030824
394          82      2      vw pickup      0.024413
395          82      1  dodge rampage      0.036601
396          82      1  ford ranger      0.030095
397          82      1  chevy s-10      0.030147

      performance_score
0          0.061644
1          0.046710
2          0.057625
3          0.055928
4          0.051754
..      ...
393          0.150968
394          0.508169
395          0.161743
396          0.198400
397          0.221103

[398 rows x 11 columns]

```

2.8 8. Visualization Preparation

Generate a summary DataFrame with: - Average mpg, weight, and horsepower for each model_year.

```
[162]: summary_df = df.groupby('model_year')[['mpg', 'weight', 'horsepower']].mean()
```

```
[163]: summary_df
```

```
[163]:
```

	mpg	weight	horsepower
model_year			
70	17.689655	3372.793103	147.827586
71	21.250000	2995.428571	106.945335
72	18.714286	3237.714286	120.178571

73	17.100000	3419.025000	130.475000
74	22.703704	2877.925926	94.609977
75	20.266667	3176.800000	101.066667
76	21.573529	3078.735294	101.117647
77	23.375000	2997.357143	105.071429
78	24.061111	2861.805556	99.694444
79	25.093103	3055.344828	101.206897
80	33.696552	2436.655172	79.342716
81	30.334483	2522.931034	81.843772
82	31.709677	2453.548387	82.208690

2.9 9. Exporting Data

Save a subset of the data containing only mpg, cylinders, horsepower, and weight for cars with mpg > 30 into a CSV file named high_mpg_cars.csv.

```
[164]: subset_df = df[df['mpg']>30]
subset_df[['mpg', 'cylinders', 'horsepower', 'weight']]
```

```
[164]:      mpg  cylinders  horsepower  weight
53   31.0           4           65.0   1773
54   35.0           4           69.0   1613
129  31.0           4           67.0   1950
131  32.0           4           65.0   1836
144  31.0           4           52.0   1649
..    ...         ...         ...    ...
390  32.0           4           96.0   2665
391  36.0           4           84.0   2370
394  44.0           4           52.0   2130
395  32.0           4           84.0   2295
397  31.0           4           82.0   2720
```

[85 rows x 4 columns]

```
[165]: subset_df[['mpg', 'cylinders', 'horsepower', 'weight']].to_csv('high_mpg_cars.csv')
```

```
[166]: verify_df = pd.read_csv('high_mpg_cars.csv')
verify_df.head()
```

```
[166]:   Unnamed: 0   mpg  cylinders  horsepower  weight
0           53  31.0           4           65.0   1773
1           54  35.0           4           69.0   1613
2          129  31.0           4           67.0   1950
3          131  32.0           4           65.0   1836
4          144  31.0           4           52.0   1649
```

2.10 10. Finding Anomalies

Identify potential outliers in the mpg column using the Interquartile Range (IQR) method. Specifically: - Calculate the IQR for mpg. - Define outliers as values less than $Q1 - 1.5 * IQR$ or greater than $Q3 + 1.5 * IQR$. - Create a DataFrame of cars classified as outliers, displaying car_name, mpg, and model_year.

```
[167]: # IQR for mpg
Q1 = df['mpg'].quantile(0.25)
Q3 = df['mpg'].quantile(0.75)
IQR = Q3-Q1
print("Q1 of mpg = ", Q1)
print("Q3 of mpg = ", Q3)
print("IQR of mpg = ", IQR)
```

```
Q1 of mpg = 17.5
Q3 of mpg = 29.0
IQR of mpg = 11.5
```

```
[168]: # Define outliers as values less than  $Q1 - 1.5 * IQR$  or greater than  $Q3 + 1.5 * IQR$ .
outliers = (df['mpg'] < (Q1 - 1.5 * IQR)) | (df['mpg'] > (Q3 + 1.5 * IQR))
print("Outliers are mpg < ", Q1 - 1.5 * IQR, "or mpg > ", Q3 + 1.5 * IQR)
```

```
Outliers are mpg < 0.25 or mpg > 46.25
```

```
[169]: # Create a DataFrame of cars classified as outliers, displaying car_name, mpg,
and model_year.
outlier_df = df[outliers]
outlier_df[['car name', 'mpg', 'model year']]
```

```
[169]:      car name    mpg  model year
322  mazda glc  46.6         80
```