

# Data\_Preprocessing\_Part\_2\_Practice

December 5, 2024

```
[51]: import numpy as np
import pandas as pd

from sklearn.preprocessing import LabelEncoder, OrdinalEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

```
[9]: data = {'id': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
            'product':
                ↳ ['apple', 'orange', 'banana', 'carrot', 'laptop', 'phone', 'shirt', 'pants', 'pen', 'box'],
            'color':
                ↳ ['red', 'orange', 'yellow', 'orange', 'black', 'white', 'blue', 'black', 'red', 'black'],
            'category':
                ↳ ['fruit', 'fruit', 'fruit', 'vegetable', 'electronic', 'electronic', 'cloth', 'cloth', 'stationary', 'stationary'],
            'sales':
                ↳ ['low', 'high', 'average', 'average', 'low', 'high', 'high', 'high', 'low', 'high'],
            'rating': [5, 4, 3, 2, 5, 4, 3, 4, 5, 4],
            'price': [100, 120, 25, 15, 80000, 20000, 300, 500, 15, 50]}

df = pd.DataFrame(data)
df
```

```
[9]:
```

	id	product	color	category	sales	rating	price
0	1	apple	red	fruit	low	5	100
1	2	orange	orange	fruit	high	4	120
2	3	banana	yellow	fruit	average	3	25
3	4	carrot	orange	vegetable	average	2	15
4	5	laptop	black	electronic	low	5	80000
5	6	phone	white	electronic	high	4	20000
6	7	shirt	blue	cloth	high	3	300
7	8	pants	black	cloth	high	4	500
8	9	pen	red	stationary	low	5	15
9	10	box	black	stationary	high	4	50

# 1 Feature Encoding

## 2 Applying label encoding on 'category' column

```
[11]: label_encoding = LabelEncoder() # create an object/instance
      # normally the original values are replaced with encoded values; new columns
      ↳ need not be created
      df['category_encoded'] = label_encoding.fit_transform(df['category'])
      df
```

```
[11]:   id product  color  category  sales  rating  price  category_encoded
0    1  apple   red    fruit     low      5    100             2
1    2  orange orange   fruit     high     4    120             2
2    3  banana yellow   fruit  average     3     25             2
3    4  carrot orange vegetable  average     2     15             4
4    5  laptop  black electronic   low      5  80000             1
5    6  phone  white electronic   high     4  20000             1
6    7  shirt   blue    cloth     high     3    300             0
7    8  pants  black    cloth     high     4    500             0
8    9   pen   red    stationary   low      5     15             3
9   10   box  black    stationary   high     4     50             3
```

```
[12]: df['product_encoded'] = label_encoding.fit_transform(df['product'])
      df
```

```
[12]:   id product  color  category  sales  rating  price  category_encoded \
0    1  apple   red    fruit     low      5    100             2
1    2  orange orange   fruit     high     4    120             2
2    3  banana yellow   fruit  average     3     25             2
3    4  carrot orange vegetable  average     2     15             4
4    5  laptop  black electronic   low      5  80000             1
5    6  phone  white electronic   high     4  20000             1
6    7  shirt   blue    cloth     high     3    300             0
7    8  pants  black    cloth     high     4    500             0
8    9   pen   red    stationary   low      5     15             3
9   10   box  black    stationary   high     4     50             3
```

```
product_encoded
0             0
1             5
2             1
3             3
4             4
5             8
6             9
7             6
8             7
```

### 3 Applying ordinal encoding on 'sales' column

```
[15]: sales_order = ['low', 'average', 'high'] # for ordinal encoding, the order of
      ↪ values must be explicitly specified
      sales_encoding = OrdinalEncoder(categories=[sales_order]) # create an object/
      ↪ instance and give the 'categories' parameter as the defined order

      df['sales_encoding'] = sales_encoding.fit_transform(df[['sales']])
      df
```

```
[15]:
```

	id	product	color	category	sales	rating	price	category_encoded	\
0	1	apple	red	fruit	low	5	100	2	
1	2	orange	orange	fruit	high	4	120	2	
2	3	banana	yellow	fruit	average	3	25	2	
3	4	carrot	orange	vegetable	average	2	15	4	
4	5	laptop	black	electronic	low	5	80000	1	
5	6	phone	white	electronic	high	4	20000	1	
6	7	shirt	blue	cloth	high	3	300	0	
7	8	pants	black	cloth	high	4	500	0	
8	9	pen	red	stationary	low	5	15	3	
9	10	box	black	stationary	high	4	50	3	

	product_encoded	sales_encoding
0	0	0.0
1	5	2.0
2	1	1.0
3	3	1.0
4	4	0.0
5	8	2.0
6	9	2.0
7	6	2.0
8	7	0.0
9	2	2.0

```
[17]: color_order = ['black', 'red', 'blue', 'yellow', 'orange', 'white'] # for ordinal
      ↪ encoding, the order of values must be explicitly specified
      color_encoding = OrdinalEncoder(categories=[color_order]) # create an object/
      ↪ instance and give the 'categories' parameter as the defined order

      df['color_encoding'] = color_encoding.fit_transform(df[['color']])
      df
```

```
[17]:
```

	id	product	color	category	sales	rating	price	category_encoded	\
0	1	apple	red	fruit	low	5	100	2	

1	2	orange	orange	fruit	high	4	120	2
2	3	banana	yellow	fruit	average	3	25	2
3	4	carrot	orange	vegetable	average	2	15	4
4	5	laptop	black	electronic	low	5	80000	1
5	6	phone	white	electronic	high	4	20000	1
6	7	shirt	blue	cloth	high	3	300	0
7	8	pants	black	cloth	high	4	500	0
8	9	pen	red	stationary	low	5	15	3
9	10	box	black	stationary	high	4	50	3

	product_encoded	sales_encoding	color_encoding
0	0	0.0	1.0
1	5	2.0	4.0
2	1	1.0	3.0
3	3	1.0	4.0
4	4	0.0	0.0
5	8	2.0	5.0
6	9	2.0	2.0
7	6	2.0	0.0
8	7	0.0	1.0
9	2	2.0	0.0

```
[ ]: # to do multiple columns at the same time
# sales_order = ['low', 'average', 'high']
# color_order = ['black', 'red', 'blue', 'yellow', 'orange', 'white']
# ordinal_encoding = OrdinalEncoder(categories=[sales_order, color_order])

# df[['sales_encoding', 'color_encoding']] = ordinal_encoding.
#   ↪ fit_transform(df[['sales', 'color']])
# df
```

## 4 On Hot Encoding

```
[18]: # Pandas method
color_onehot_encoding = pd.get_dummies(df['color'])
color_onehot_encoding
```

```
[18]:   black  blue  orange  red  white  yellow
0  False  False  False  True  False  False
1  False  False   True  False  False  False
2  False  False  False  False  False  True
3  False  False   True  False  False  False
4   True  False  False  False  False  False
5  False  False  False  False  True  False
6  False  True  False  False  False  False
7   True  False  False  False  False  False
```

```

8 False False False True False False
9  True False False False False False

```

```

[19]: color_onehot_encoding = pd.get_dummies(df,columns=['color'])
color_onehot_encoding

```

```

[19]:
  id product  category  sales  rating  price  category_encoded  \
0  1  apple    fruit    low     5    100                2
1  2  orange    fruit    high    4    120                2
2  3  banana    fruit  average    3     25                2
3  4  carrot  vegetable  average    2     15                4
4  5  laptop  electronic    low     5  80000                1
5  6  phone  electronic    high    4  20000                1
6  7  shirt    cloth    high     3     300                0
7  8  pants    cloth    high     4     500                0
8  9    pen  stationary    low     5     15                3
9 10    box  stationary    high    4     50                3

  product_encoded  sales_encoding  color_encoding  color_black  color_blue  \
0                0             0.0              1.0         False        False
1                5             2.0              4.0         False        False
2                1             1.0              3.0         False        False
3                3             1.0              4.0         False        False
4                4             0.0              0.0          True        False
5                8             2.0              5.0         False        False
6                9             2.0              2.0         False         True
7                6             2.0              0.0          True        False
8                7             0.0              1.0         False        False
9                2             2.0              0.0          True        False

  color_orange  color_red  color_white  color_yellow
0          False        True        False        False
1           True        False        False        False
2          False        False        False         True
3           True        False        False        False
4          False        False        False        False
5          False        False         True        False
6          False        False        False        False
7          False        False        False        False
8          False         True        False        False
9          False        False        False        False

```

```

[43]: #can change prefix, dtype of new columns and also specify drop_first as True to
      ↪avoid multi-collinearity (dummy variable trap)
color_onehot_encoding = pd.
      ↪get_dummies(df,columns=['color'],prefix='col',dtype='int', drop_first=True)
color_onehot_encoding

```

```
[43]:
```

	id	product	category	sales	rating	price	category_encoded	\
0	1	apple	fruit	low	5	100	2	
1	2	orange	fruit	high	4	120	2	
2	3	banana	fruit	average	3	25	2	
3	4	carrot	vegetable	average	2	15	4	
4	5	laptop	electronic	low	5	80000	1	
5	6	phone	electronic	high	4	20000	1	
6	7	shirt	cloth	high	3	300	0	
7	8	pants	cloth	high	4	500	0	
8	9	pen	stationary	low	5	15	3	
9	10	box	stationary	high	4	50	3	

  

	product_encoded	sales_encoding	color_encoding	col_blue	col_orange	\
0	0	0.0	1.0	0	0	
1	5	2.0	4.0	0	1	
2	1	1.0	3.0	0	0	
3	3	1.0	4.0	0	1	
4	4	0.0	0.0	0	0	
5	8	2.0	5.0	0	0	
6	9	2.0	2.0	1	0	
7	6	2.0	0.0	0	0	
8	7	0.0	1.0	0	0	
9	2	2.0	0.0	0	0	

  

	col_red	col_white	col_yellow
0	1	0	0
1	0	0	0
2	0	0	1
3	0	0	0
4	0	0	0
5	0	1	0
6	0	0	0
7	0	0	0
8	1	0	0
9	0	0	0

```
[49]: # using sklearn
one_hot_encoding = OneHotEncoder()
color_mod_encoding = one_hot_encoding.fit_transform(df[['color']])
print(color_mod_encoding)
```

```
(0, 3)      1.0
(1, 2)      1.0
(2, 5)      1.0
(3, 2)      1.0
(4, 0)      1.0
(5, 4)      1.0
(6, 1)      1.0
```

```
(7, 0)          1.0
(8, 3)          1.0
(9, 0)          1.0
```

```
[48]: one_hot_df = pd.DataFrame(color_mod_encoding, columns=one_hot_encoding.
    ↪get_feature_names_out(['color']))
# OR
# one_hot_df = pd.DataFrame(color_mod_encoding,
    ↪columns=['black', 'red', 'blue', 'yellow', 'orange', 'white'])
one_hot_df
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[48], line 1
----> 1 one_hot_df = pd.DataFrame(color_mod_encoding, columns=one_hot_encoding.
    ↪get_feature_names_out(['color']))
      2 # OR
      3 one_hot_df

File /opt/anaconda3/lib/python3.12/site-packages/pandas/core/frame.py:867, in
    ↪DataFrame.__init__(self, data, index, columns, dtype, copy)
      859         mgr = arrays_to_mgr(
      860             arrays,
      861             columns,
      (...),
      864             typ=manager,
      865         )
      866     else:
--> 867         mgr = ndarray_to_mgr(
      868             data,
      869             index,
      870             columns,
      871             dtype=dtype,
      872             copy=copy,
      873             typ=manager,
      874         )
      875 else:
      876     mgr = dict_to_mgr(
      877         {},
      878         index,
      (...),
      881         typ=manager,
      882     )

File /opt/anaconda3/lib/python3.12/site-packages/pandas/core/internals/
    ↪construction.py:336, in ndarray_to_mgr(values, index, columns, dtype, copy,
    ↪typ)
      331 # _prep_ndarraylike ensures that values.ndim == 2 at this point
```

```

332 index, columns = _get_axes(
333     values.shape[0], values.shape[1], index=index, columns=columns
334 )
--> 336 _check_values_indices_shape_match(values, index, columns)
338 if typ == "array":
339     if issubclass(values.dtype.type, str):

File /opt/anaconda3/lib/python3.12/site-packages/pandas/core/internals/
↳ construction.py:420, in _check_values_indices_shape_match(values, index,
↳ columns)
    418 passed = values.shape
    419 implied = (len(index), len(columns))
--> 420 raise ValueError(f"Shape of passed values is {passed}, indices imply
↳ {implied}")

ValueError: Shape of passed values is (10, 1), indices imply (10, 6)

```

```
[ ]: new_df = pd.concat(df,one_hot_df,axis=1)
new_df
```

## 5 beer\_servings.csv

```
[26]: # continent - one hot
df_beer = pd.read_csv('beer-servings.csv')
df_beer
```

```
[26]:
```

	Unnamed: 0	country	beer_servings	spirit_servings	wine_servings	\
0	0	Afghanistan	0.0	0.0	0.0	
1	1	Albania	89.0	132.0	54.0	
2	2	Algeria	25.0	0.0	14.0	
3	3	Andorra	245.0	138.0	312.0	
4	4	Angola	217.0	57.0	45.0	
..	...	...	...	...	...	
188	188	Venezuela	NaN	100.0	3.0	
189	189	Vietnam	111.0	2.0	1.0	
190	190	Yemen	6.0	0.0	0.0	
191	191	Zambia	32.0	19.0	4.0	
192	192	Zimbabwe	64.0	18.0	4.0	

  

	total_litres_of_pure_alcohol	continent
0	0.0	Asia
1	4.9	Europe
2	0.7	Africa
3	12.4	Europe
4	5.9	Africa
..	...	...



```

188             7.7  South America
189             2.0             Asia
190             0.1             Asia
191             2.5             Africa
192             4.7             Africa

```

[193 rows x 7 columns]

```

[28]: # country - label encoding
df_beer['country_encoding'] = label_encoding.fit_transform(df_beer['country'])
df_beer

```

```

[28]:      Unnamed: 0      country  beer_servings  spirit_servings  wine_servings  \
0             0  Afghanistan             0.0             0.0             0.0
1             1    Albania             89.0            132.0             54.0
2             2    Algeria             25.0              0.0             14.0
3             3    Andorra            245.0            138.0            312.0
4             4    Angola             217.0             57.0             45.0
..          ...          ...          ...          ...          ...
188          188    Venezuela             NaN            100.0              3.0
189          189    Vietnam            111.0              2.0              1.0
190          190     Yemen              6.0              0.0              0.0
191          191    Zambia              32.0             19.0              4.0
192          192    Zimbabwe             64.0             18.0              4.0

      total_litres_of_pure_alcohol  continent  country_encoding
0              0.0             Asia              0
1              4.9             Europe              1
2              0.7             Africa              2
3             12.4             Europe              3
4              5.9             Africa              4
..          ...          ...          ...
188             7.7  South America            188
189             2.0             Asia            189
190             0.1             Asia            190
191             2.5             Africa            191
192             4.7             Africa            192

```

[193 rows x 8 columns]

```

[42]: # continent - one hot encoding
continent_onehot_encoding = pd.
↳ get_dummies(df_beer, columns=['continent'], prefix='cont', dtype='int', drop_first=True)
continent_onehot_encoding

```

```

[42]:      Unnamed: 0      country  beer_servings  spirit_servings  wine_servings  \
0             0  Afghanistan             0.0             0.0             0.0

```

1	1	Albania	89.0	132.0	54.0
2	2	Algeria	25.0	0.0	14.0
3	3	Andorra	245.0	138.0	312.0
4	4	Angola	217.0	57.0	45.0
..	...	...	...	...	...
188	188	Venezuela	NaN	100.0	3.0
189	189	Vietnam	111.0	2.0	1.0
190	190	Yemen	6.0	0.0	0.0
191	191	Zambia	32.0	19.0	4.0
192	192	Zimbabwe	64.0	18.0	4.0

	total_litres_of_pure_alcohol	country_encoding	cont_Asia	cont_Europe	\
0	0.0	0	1	0	
1	4.9	1	0	1	
2	0.7	2	0	0	
3	12.4	3	0	1	
4	5.9	4	0	0	
..	...	...	...	...	
188	7.7	188	0	0	
189	2.0	189	1	0	
190	0.1	190	1	0	
191	2.5	191	0	0	
192	4.7	192	0	0	

	cont_North America	cont_Oceania	cont_South America
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0
..	...	...	...
188	0	0	1
189	0	0	0
190	0	0	0
191	0	0	0
192	0	0	0

[193 rows x 12 columns]

## 6 Scaling

## 7 Standard Scaling

```
[53]: std_scaler = StandardScaler() # instance
df['price_standardised_scale'] = std_scaler.fit_transform(df[['price']])
df
```

```
[53]:
```

	id	product	color	category	sales	rating	price	category_encoded	\
0	1	apple	red	fruit	low	5	100	2	
1	2	orange	orange	fruit	high	4	120	2	
2	3	banana	yellow	fruit	average	3	25	2	
3	4	carrot	orange	vegetable	average	2	15	4	
4	5	laptop	black	electronic	low	5	80000	1	
5	6	phone	white	electronic	high	4	20000	1	
6	7	shirt	blue	cloth	high	3	300	0	
7	8	pants	black	cloth	high	4	500	0	
8	9	pen	red	stationary	low	5	15	3	
9	10	box	black	stationary	high	4	50	3	

  

	product_encoded	sales_encoding	color_encoding	price_standardised_scale
0	0	0.0	1.0	-0.416546
1	5	2.0	4.0	-0.415714
2	1	1.0	3.0	-0.419667
3	3	1.0	4.0	-0.420083
4	4	0.0	0.0	2.907505
5	8	2.0	5.0	0.411346
6	9	2.0	2.0	-0.408226
7	6	2.0	0.0	-0.399905
8	7	0.0	1.0	-0.420083
9	2	2.0	0.0	-0.418627

```
[54]: df['price'].describe()
```

```
[54]: count      10.000000
mean      10112.500000
std       25337.151933
min        15.000000
25%        31.250000
50%       110.000000
75%       450.000000
max      80000.000000
Name: price, dtype: float64
```

```
[56]: df['price_standardised_scale'].describe()
```

```
[56]: count    1.000000e+01
      mean     2.220446e-17
      std      1.054093e+00
      min     -4.200827e-01
      25%     -4.194067e-01
      50%     -4.161305e-01
      75%     -4.019856e-01
      max      2.907505e+00
      Name: price_standardised_scale, dtype: float64
```

## 8 Min-Max Scaling

```
[57]: norm_scaler = MinMaxScaler() # instance
      # norm_scaler = MinMaxScaler(feature_reange=(0,1)) # instance
      df['price_normalised_scale'] = norm_scaler.fit_transform(df[['price']])
      df
```

```
[57]:
```

	id	product	color	category	sales	rating	price	category_encoded	\
0	1	apple	red	fruit	low	5	100	2	
1	2	orange	orange	fruit	high	4	120	2	
2	3	banana	yellow	fruit	average	3	25	2	
3	4	carrot	orange	vegetable	average	2	15	4	
4	5	laptop	black	electronic	low	5	80000	1	
5	6	phone	white	electronic	high	4	20000	1	
6	7	shirt	blue	cloth	high	3	300	0	
7	8	pants	black	cloth	high	4	500	0	
8	9	pen	red	stationary	low	5	15	3	
9	10	box	black	stationary	high	4	50	3	

  

	product_encoded	sales_encoding	color_encoding	price_standardised_scale	\
0	0	0.0	1.0	-0.416546	
1	5	2.0	4.0	-0.415714	
2	1	1.0	3.0	-0.419667	
3	3	1.0	4.0	-0.420083	
4	4	0.0	0.0	2.907505	
5	8	2.0	5.0	0.411346	
6	9	2.0	2.0	-0.408226	
7	6	2.0	0.0	-0.399905	
8	7	0.0	1.0	-0.420083	
9	2	2.0	0.0	-0.418627	

  

	price_normalised_scale
0	0.001063
1	0.001313
2	0.000125
3	0.000000

4	1.000000
5	0.249859
6	0.003563
7	0.006064
8	0.000000
9	0.000438

```
[58]: df['price'].describe()
```

```
[58]: count      10.000000
      mean      10112.500000
      std       25337.151933
      min       15.000000
      25%       31.250000
      50%      110.000000
      75%      450.000000
      max      80000.000000
      Name: price, dtype: float64
```

```
[60]: df['price_normalised_scale'].describe()
```

```
[60]: count      10.000000
      mean       0.126242
      std       0.316774
      min       0.000000
      25%       0.000203
      50%       0.001188
      75%       0.005439
      max       1.000000
      Name: price_normalised_scale, dtype: float64
```

## 9 Correlation check

```
[67]: # Read dataset
      df_beer = pd.read_csv('beer-servings.csv')

      # Deleting "Unnamed:0" column because it is useless
      df_beer = df_beer.iloc[:,1:]

      # to delete any duplicate rows
      df_beer.drop_duplicates(inplace=True)

      # Splitting numerical and categorical columns
      num_df_beer = df_beer.select_dtypes(include="number")
      cat_df_beer = df_beer.select_dtypes(include="object_")

      num_cols = num_df_beer.columns.tolist()
```

```

cat_cols = cat_df_beer.columns.tolist()

# Filling missing values for numerical columns
for col in num_cols:
    df_beer[col] = df_beer[col].fillna(df_beer[col].median())

# Filling missing values for categorical columns
for col in cat_cols:
    df_beer[col] = df_beer[col].fillna(df_beer[col].mode()[0])

def remove_outliers(df, column_name):
    q1 = df[column_name].quantile(0.25)
    q3 = df[column_name].quantile(0.75)
    iqr = q3-q1
    lower_bound = q1 - 1.5*iqr
    upper_bound = q3 + 1.5*iqr
    df[column_name] = df[column_name].clip(upper=upper_bound)
    df[column_name] = df[column_name].clip(lower=lower_bound)
    return df[column_name]

for col in num_cols:
    df_beer[col] = remove_outliers(df_beer, col)

df_beer

```

```

[67]:
      country  beer_servings  spirit_servings  wine_servings \
0   Afghanistan          0.0             0.0             0.0
1     Albania          89.0            132.0            54.0
2     Algeria          25.0             0.0            14.0
3    Andorra          245.0            138.0           146.0
4     Angola          217.0             57.0            45.0
..         ...            ...            ...            ...
188  Venezuela          76.0            100.0             3.0
189   Vietnam          111.0             2.0             1.0
190     Yemen           6.0             0.0             0.0
191    Zambia          32.0             19.0             4.0
192   Zimbabwe          64.0             18.0             4.0

      total_litres_of_pure_alcohol  continent
0                                0.0        Asia
1                                4.9        Europe
2                                0.7        Africa
3                               12.4        Europe
4                                5.9        Africa
..                                ...            ...
188                             7.7  South America
189                             2.0            Asia

```

190	0.1	Asia
191	2.5	Africa
192	4.7	Africa

[193 rows x 6 columns]

```
[68]: # Make copies of the dataset to apply std and norm scaling
df_beer_std_copy = df_beer
df_beer_norm_copy = df_beer
```

```
[74]: std_df = std_scaler.fit_transform(df_beer_std_copy[num_cols])
pd.DataFrame(std_df)
```

```
[74]:
```

	0	1	2	3
0	-1.056880	-0.971267	-0.723136	-1.264356
1	-0.151713	0.608759	0.302520	0.042922
2	-0.802619	-0.971267	-0.457225	-1.077602
3	1.434871	0.680578	2.049934	2.043856
4	1.150100	-0.288983	0.131578	0.309713
..	...	...	...	...
188	-0.283928	0.225722	-0.666155	0.789937
189	0.072036	-0.947327	-0.704142	-0.730773
190	-0.995857	-0.971267	-0.723136	-1.237677
191	-0.731427	-0.743839	-0.647161	-0.597378
192	-0.405973	-0.755809	-0.647161	-0.010437

[193 rows x 4 columns]

```
[75]: norm_df = norm_scaler.fit_transform(df_beer_norm_copy[num_cols])
pd.DataFrame(norm_df)
```

```
[75]:
```

	0	1	2	3
0	0.000000	0.0000	0.000000	0.000000
1	0.236702	0.4224	0.369863	0.340278
2	0.066489	0.0000	0.095890	0.048611
3	0.651596	0.4416	1.000000	0.861111
4	0.577128	0.1824	0.308219	0.409722
..	...	...	...	...
188	0.202128	0.3200	0.020548	0.534722
189	0.295213	0.0064	0.006849	0.138889
190	0.015957	0.0000	0.000000	0.006944
191	0.085106	0.0608	0.027397	0.173611
192	0.170213	0.0576	0.027397	0.326389

[193 rows x 4 columns]

```
[79]: # Correlation should not be affected even after all the preprocessing is done
print("Correlation of original df = \n", pd.DataFrame(df_beer[num_cols]).corr())

print("\n\nCorrelation of standardised df = \n", pd.DataFrame(std_df).corr())

print("\n\nCorrelation of normalised df = \n", pd.DataFrame(norm_df).corr())
```

Correlation of original df =

	beer_servings	spirit_servings	wine_servings	\
beer_servings	1.000000	0.473518	0.617509	
spirit_servings	0.473518	1.000000	0.280068	
wine_servings	0.617509	0.280068	1.000000	
total_litres_of_pure_alcohol	0.829418	0.659014	0.716568	

	total_litres_of_pure_alcohol
beer_servings	0.829418
spirit_servings	0.659014
wine_servings	0.716568
total_litres_of_pure_alcohol	1.000000

Correlation of standardised df =

	0	1	2	3
0	1.000000	0.473518	0.617509	0.829418
1	0.473518	1.000000	0.280068	0.659014
2	0.617509	0.280068	1.000000	0.716568
3	0.829418	0.659014	0.716568	1.000000

Correlation of normalised df =

	0	1	2	3
0	1.000000	0.473518	0.617509	0.829418
1	0.473518	1.000000	0.280068	0.659014
2	0.617509	0.280068	1.000000	0.716568
3	0.829418	0.659014	0.716568	1.000000

[ ]:

[ ]: